Meeting Notes — NLB Document Intelligence: Project Kickoff

Date: 2024-03-11 **Location:** NLB HQ, Trg republike, Ljubljana (hybrid) **Kvant attendees:** Tine Hribar, Maša Vidmar, Anja Štrukelj **NLB attendees:** Petra Kobal (Head of IT Transformation), Damjan Repič (Compliance), Urška Fink (Legal Operations) **Note-taker:** Anja Štrukelj

---

Context

NLB reached out in January after seeing Tine's talk at the Digital Slovenia Forum. Their stated problem: legal and compliance teams spend roughly 3–4 hours per analyst per day manually searching through internal policy documents, regulatory circulars (Bank of Slovenia), and contract archives. The ask is a conversational search interface that can handle natural language queries in Slovenian and return answers with source citations.

This is the first formal meeting after the scoping call on Feb 22.

Discussion

Scope confirmation

Petra confirmed the corpus is approximately 4,800 documents: - Internal policy manuals (~600 docs, updated quarterly) - Bank of Slovenia regulatory circulars, 2018–present (~1,200 docs) - Contract templates and signed agreements (~2,400 docs — **these are sensitive, need separate access tier**) - Meeting minutes from compliance committee (~600 docs)

Damjan flagged that the regulatory circulars sometimes contradict each other when regulations are updated — older circulars are not always removed from the repository. This is a known pain point. Their current Confluence search just returns everything and lets users figure it out.

Urška asked whether the system could flag when a document has been superseded. We said yes in principle but that this requires metadata we may need to extract rather than rely on NLB to provide. She seemed skeptical about how reliable that would be. Fair concern, honestly — we'll need to think about this.

Technical discussion

Maša walked through the proposed architecture: - Chunking strategy TBD (policy docs are long and structured; contracts have boilerplate that will confuse retrieval if not handled carefully) - Embeddings: we're planning `text-embedding-3-large` as baseline, may test Voyage for Slovenian-language

performance - Vector DB: Qdrant hosted on Hetzner (NLB has a preference for EU data residency — AWS Frankfurt or Hetzner are both acceptable) - LLM for generation: GPT-4o initially, with the option to switch to a locally-hosted model if compliance demands it (Damjan raised this and we said we'd explore Llama 3.1 70B on dedicated hardware, but made no commitment)

Petra asked about hallucination risk. Maša explained the citation approach and the "I don't know" handling. Petra visibly relaxed a bit. Damjan did not.

Access and data transfer

NLB will not give us raw document access. They want to set up a secure transfer protocol — likely SFTP to a Hetzner instance that we provision, with NLB's security team reviewing the setup. Damjan said this could take 3–4 weeks minimum. That's going to slip the timeline. We did not argue about this in the meeting.

The contract archive is explicitly out of scope for phase 1. Urška pushed back briefly but Damjan shut it down. Internal NLB politics — not our problem.

Decisions

- Phase 1 corpus: policy manuals + regulatory circulars only (~1,800 docs)
- Language: Slovenian-first, English fallback
- Data residency: Hetzner Falkenstein (Germany)
- Timeline: delayed until data access is resolved — realistic start on ingestion pipeline is **April 7 at the earliest**
- Next meeting: March 25, agenda is data transfer protocol + metadata schema

---

Open Questions

#	Question	Owner	Due
1	Can NLB provide document-level metadata (date, author, status: active/superseded)?	Petra Kobal	March 22
2	What is the acceptable latency for query responses? (Damjan mentioned <5s as a hard requirement)	Anja to confirm in writing	March 15

#	Question	Owner	Due
3	Do we need SSO integration with NLB's Active Directory for the UI?	Petra Kobal	March 25
4	Llama 3.1 70B on-prem — is this worth scoping as an option?	Maša Vidmar	Internal, March 20

Internal Notes (not shared with client)

Damjan Repič is going to be a problem. He's risk-averse in the way that makes projects expand in scope without the budget following. The "no hallucinations" expectation needs to be managed in writing, not just verbally — Anja will draft a short explainer doc on LLM limitations for NLB's compliance team before the next meeting.

The 4–6 week data access delay means we can use that time to prototype on synthetic/public data. Tine wants to show something at the March 25 meeting — possible but it'll be shallow.

Budget confirmed at €48,000 for phase 1. No phase 2 commitment yet.