Proposal: Mercator Customer Support Chatbot (Phase 2 Follow-Up)

Prepared by: Kvant Consulting **For:** Mercator d.d. — Customer Experience Division **Date:** 2024-10-03 **Version:** 1.0 **Contact:** Anja Štrukelj (anja@kvant.si)

Context

Following the successful internal deployment of the Polica knowledge assistant (June 2024), Mercator's CX team has asked Kvant to scope a customer-facing chatbot for their webshop and loyalty app. The ask is narrow: handle the top 40 support query types (order status, return policy, loyalty points, store hours) without routing to a human agent. Volume is approximately 3,200 contacts/day across chat and email.

This is a follow-up engagement, not a new client relationship. We know the Mercator IT environment.

Proposed Solution

A retrieval-augmented chatbot grounded on Mercator's support documentation (~180 articles), returns/exchange policy, and live order data via API. Claude Haiku for generation (latency-sensitive; customer-facing). Fallback to human agent if confidence below threshold or query type outside scope.

Key design constraint: no hallucination on return policy. If the answer isn't in the docs, the bot must say so and transfer. This is non-negotiable and will be enforced at the retrieval confidence layer.

Scope

In scope: - RAG pipeline over support docs (Slovenian) - Integration with Mercator's order status API (read-only) - Escalation logic to human agent queue - Basic analytics dashboard (query volume, escalation rate, top query types)

Out of scope: - Email channel (phase 2 if this works) - Personalization based on purchase history - Any write operations to Mercator systems - Mobile app integration (separate project)

Timeline

Phase	Duration	Deliverable
Discovery + doc audit	Week 1–2	Corpus assessment, API spec review
Pipeline + integration build	Week 3–6	Working prototype, internal testing
UAT with Mercator CX team	Week 7–8	Bug fixes, threshold tuning
Soft launch (10% traffic)	Week 9	Live monitoring, escalation rate tracking
Full launch + handover	Week 10	Deployment docs, runbook

Team

Anja Štrukelj (lead), Rok Lampret (engineering), Luka Zorko (integration support).

Pricing

Fixed fee: €18,500

Includes all development, testing, and one month post-launch support. Infrastructure (~€60/month) transfers to Mercator after handover. OpenAI API costs not included — estimated €200–350/month at projected volume.

Assumptions and Risks

- Mercator's order status API is documented and accessible; undocumented endpoints will require additional scoping.
- Support documentation is current and maintained. If articles are outdated, retrieval quality will suffer — a doc audit in week 1 will surface this.
- Slovenian-language performance of Claude Haiku is acceptable; we have tested this on Polica-adjacent content and are confident, but customer-facing queries are shorter and more colloquial, which may behave differently.
- Timeline assumes Mercator CX team provides 2 reviewers for UAT week 7–8.